

William Stallings

Computer Organization

and Architecture

8th Edition

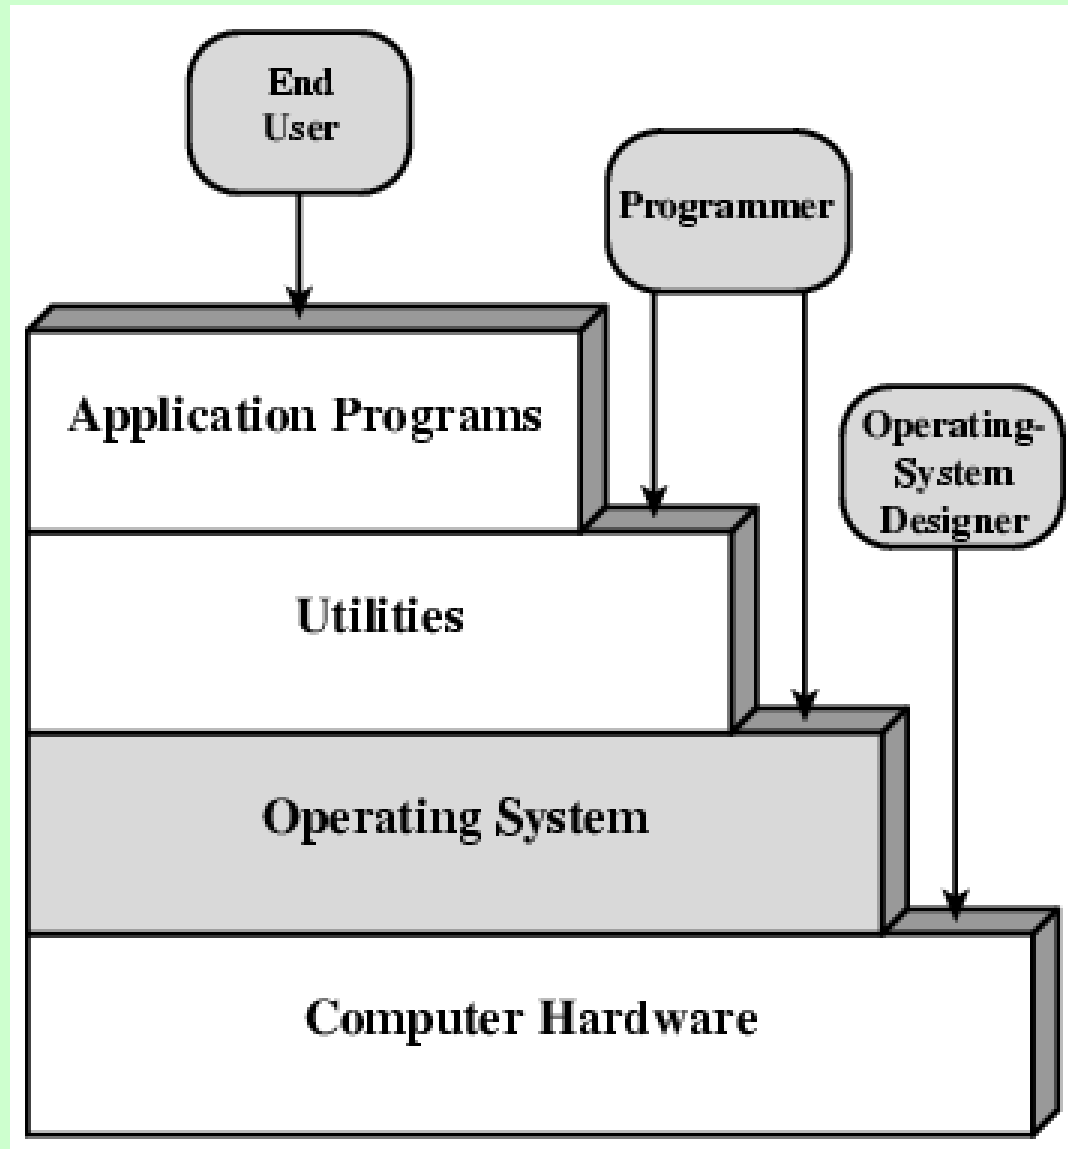
Chapter 8

Operating System Support

Objectives and Functions

- Convenience
 - Making the computer easier to use
- Efficiency
 - Allowing better use of computer resources

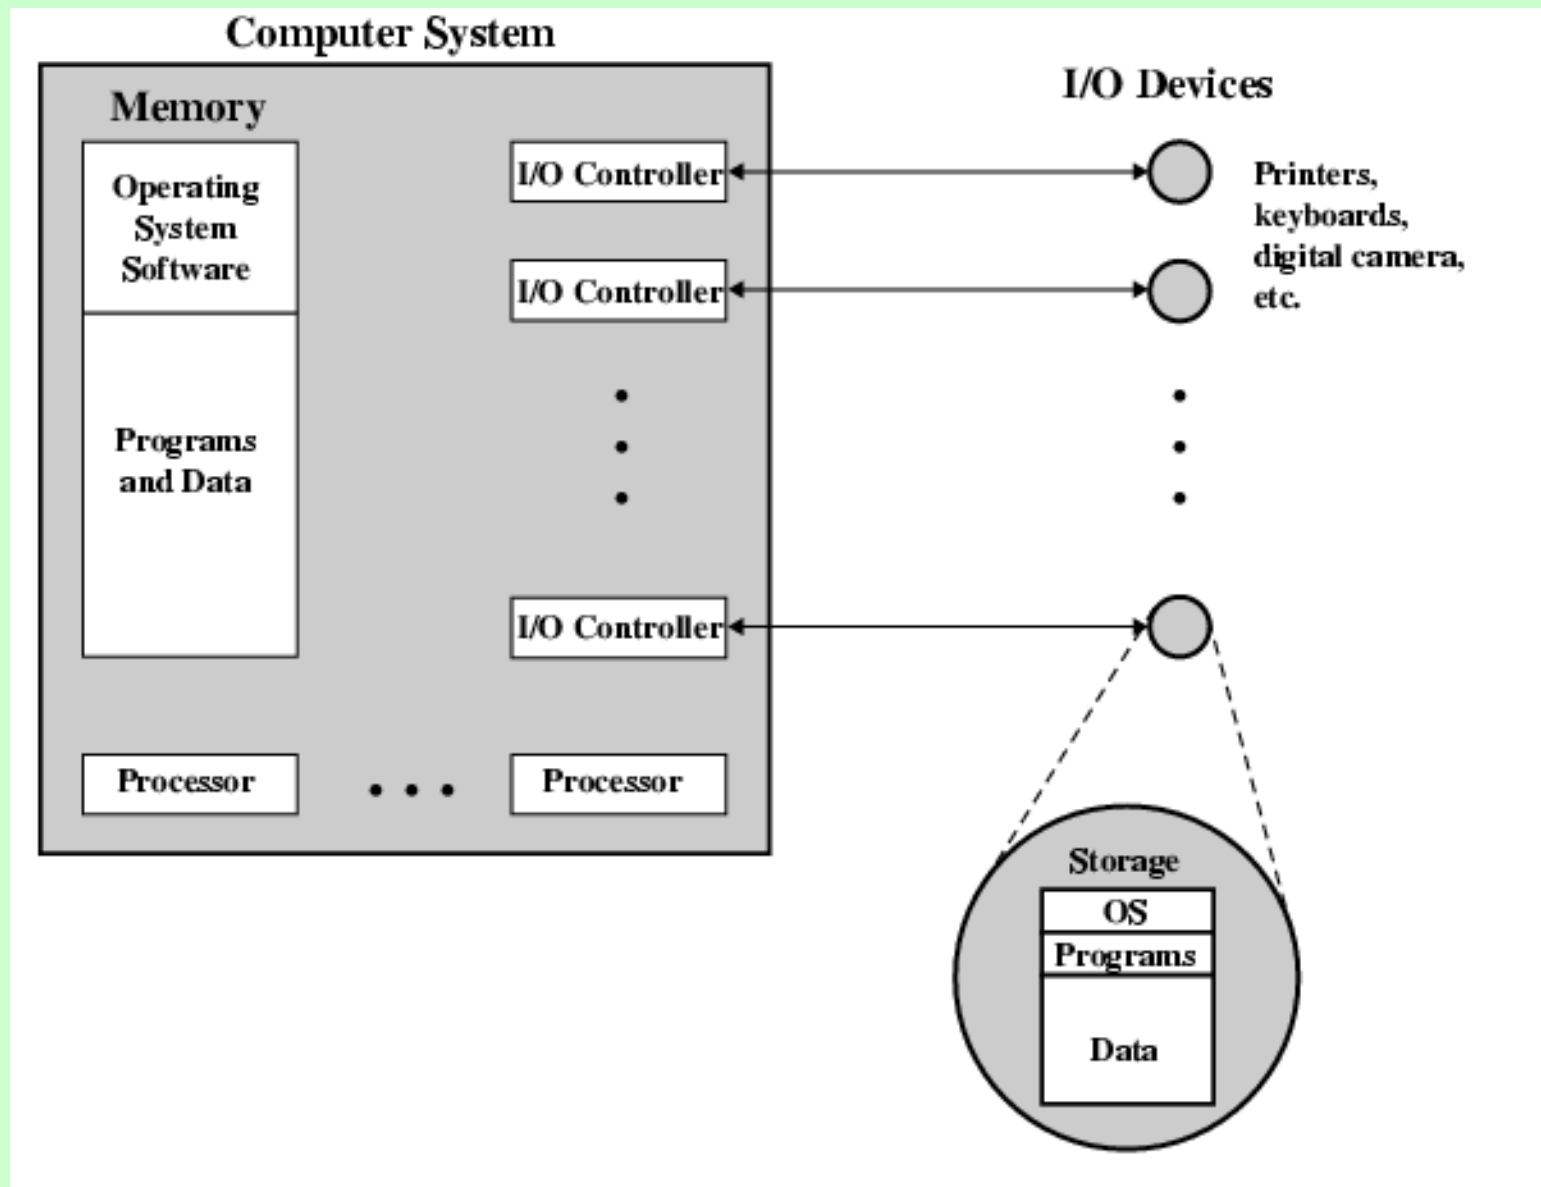
Layers and Views of a Computer System



Operating System Services

- Program creation
- Program execution
- Access to I/O devices
- Controlled access to files
- System access
- Error detection and response
- Accounting

O/S as a Resource Manager



Types of Operating System

- Interactive
- Batch
- Single program (Uni-programming)
- Multi-programming (Multi-tasking)

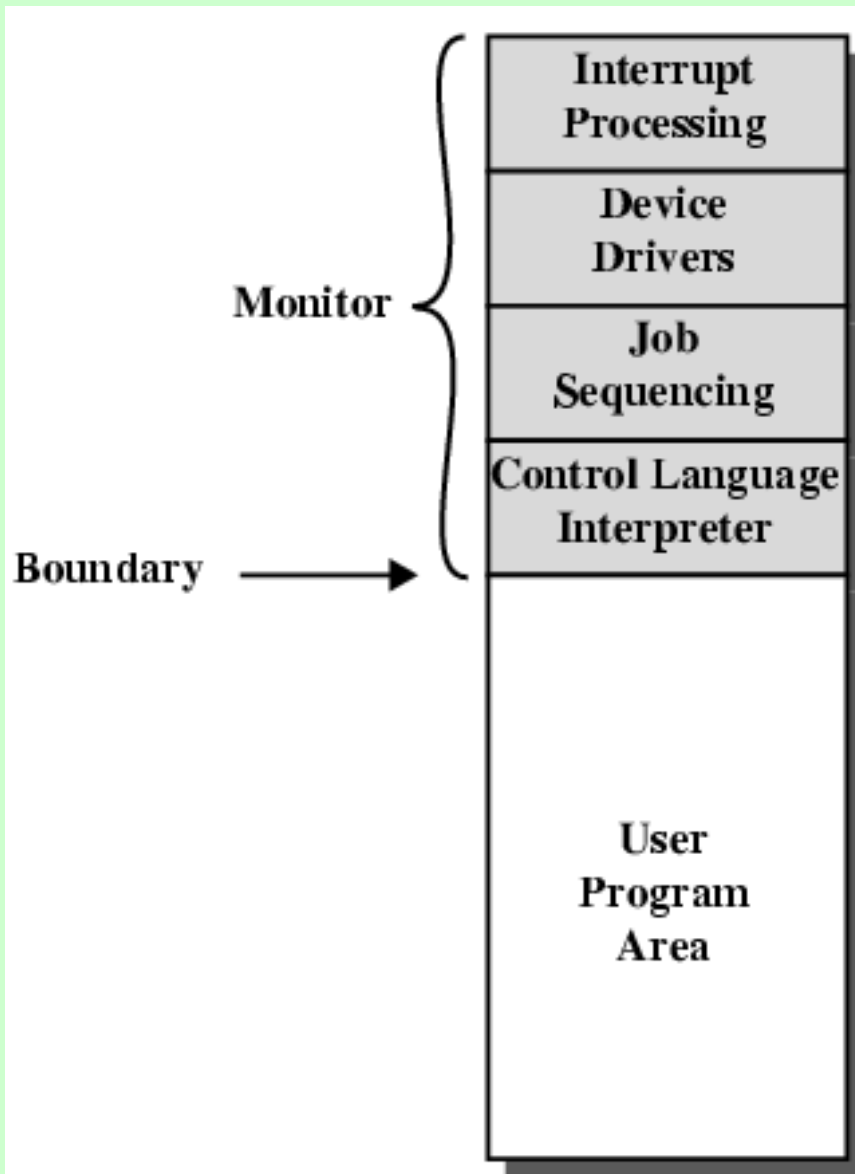
Early Systems

- Late 1940s to mid 1950s
- No Operating System
- Programs interact directly with hardware
- Two main problems:
 - Scheduling
 - Setup time

Simple Batch Systems

- Resident Monitor program
- Users submit jobs to operator
- Operator batches jobs
- Monitor controls sequence of events to process batch
- When one job is finished, control returns to Monitor which reads next job
- Monitor handles scheduling

Memory Layout for Resident Monitor



Job Control Language

- Instructions to Monitor
- Usually denoted by \$
- e.g.
 - \$JOB
 - \$FTN
 - ... Some Fortran instructions
 - \$LOAD
 - \$RUN
 - ... Some data
 - \$END

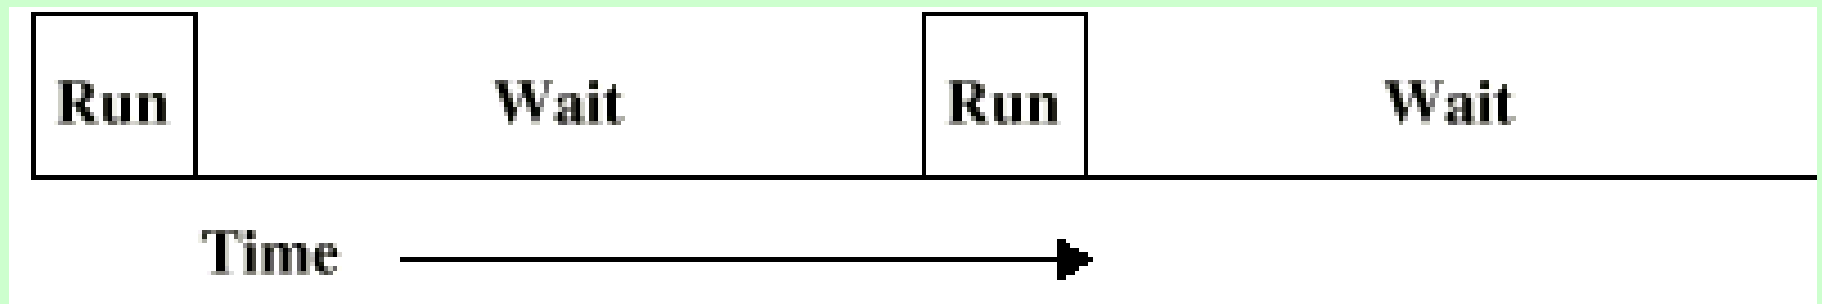
Desirable Hardware Features

- Memory protection
 - To protect the Monitor
- Timer
 - To prevent a job monopolizing the system
- Privileged instructions
 - Only executed by Monitor
 - e.g. I/O
- Interrupts
 - Allows for relinquishing and regaining control

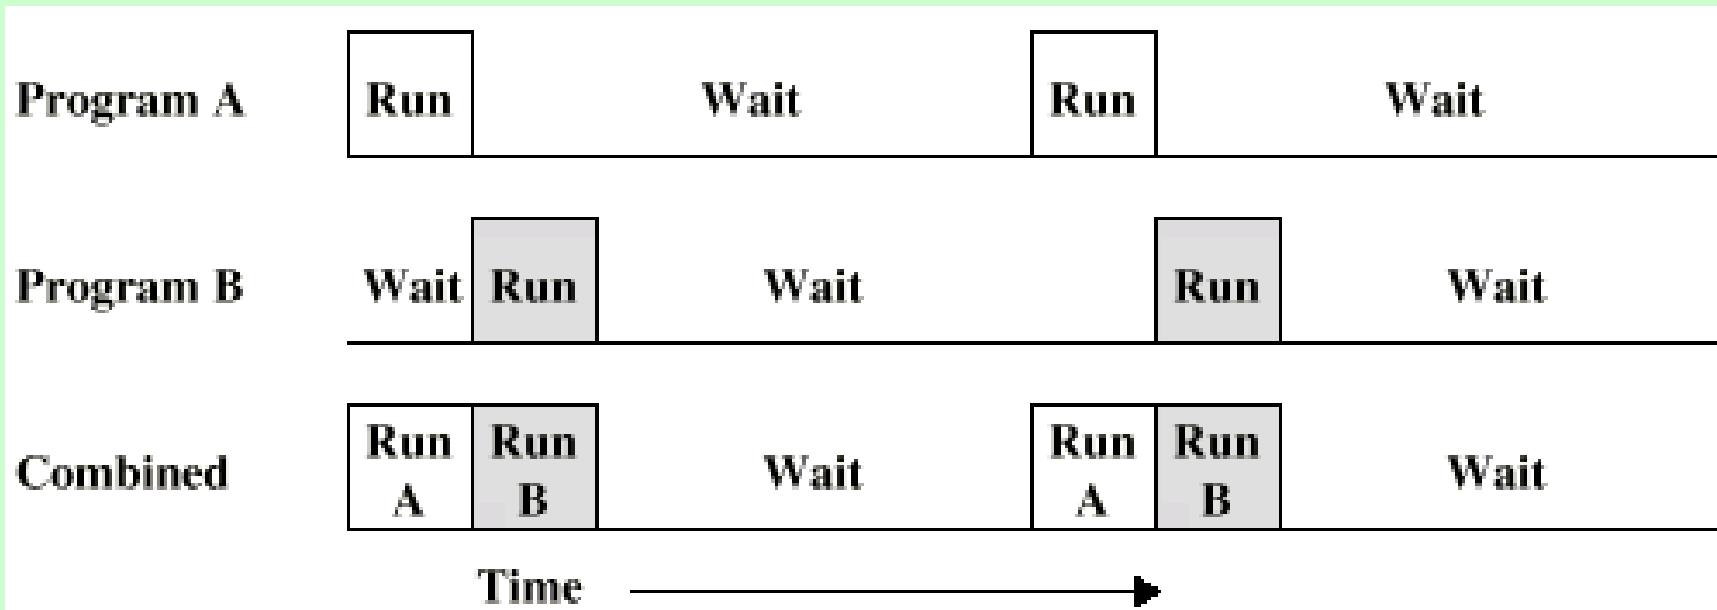
Multi-programmed Batch Systems

- I/O devices very slow
- When one program is waiting for I/O, another can use the CPU

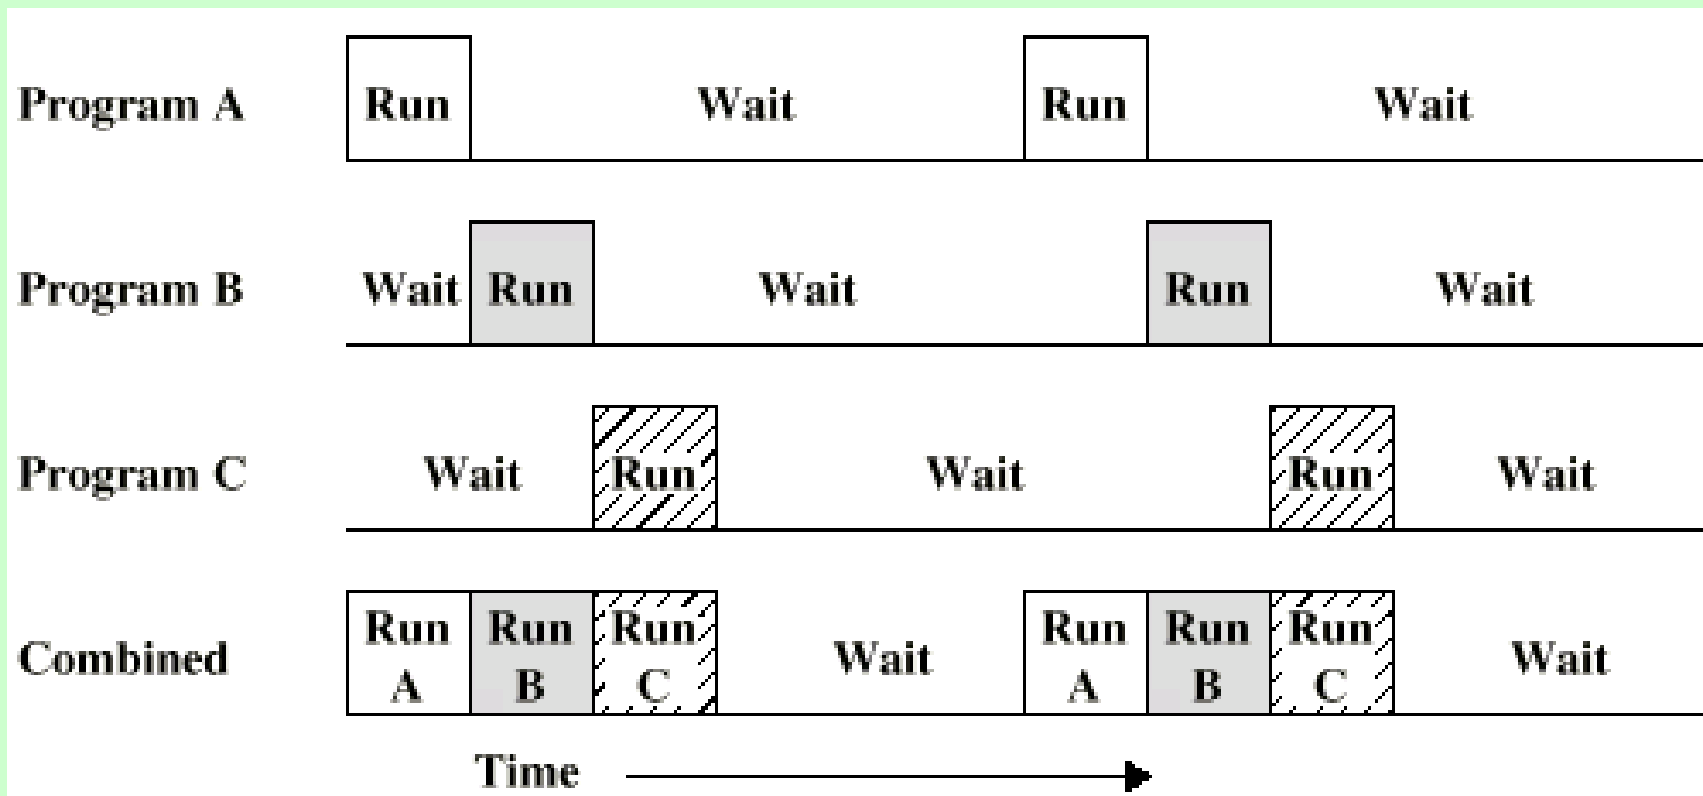
Single Program




Multi-Programming with Two Programs




Multi-Programming with Three Programs



Utilization



(a) Uniprogramming



(b) Multiprogramming

Time Sharing Systems

- Allow users to interact directly with the computer
 - i.e. Interactive
- Multi-programming allows a number of users to interact with the computer

Scheduling

- Key to multi-programming
- Long term
- Medium term
- Short term
- I/O

Long Term Scheduling

- Determines which programs are submitted for processing
- i.e. controls the degree of multi-programming
- Once submitted, a job becomes a process for the short term scheduler
- (or it becomes a swapped out job for the medium term scheduler)

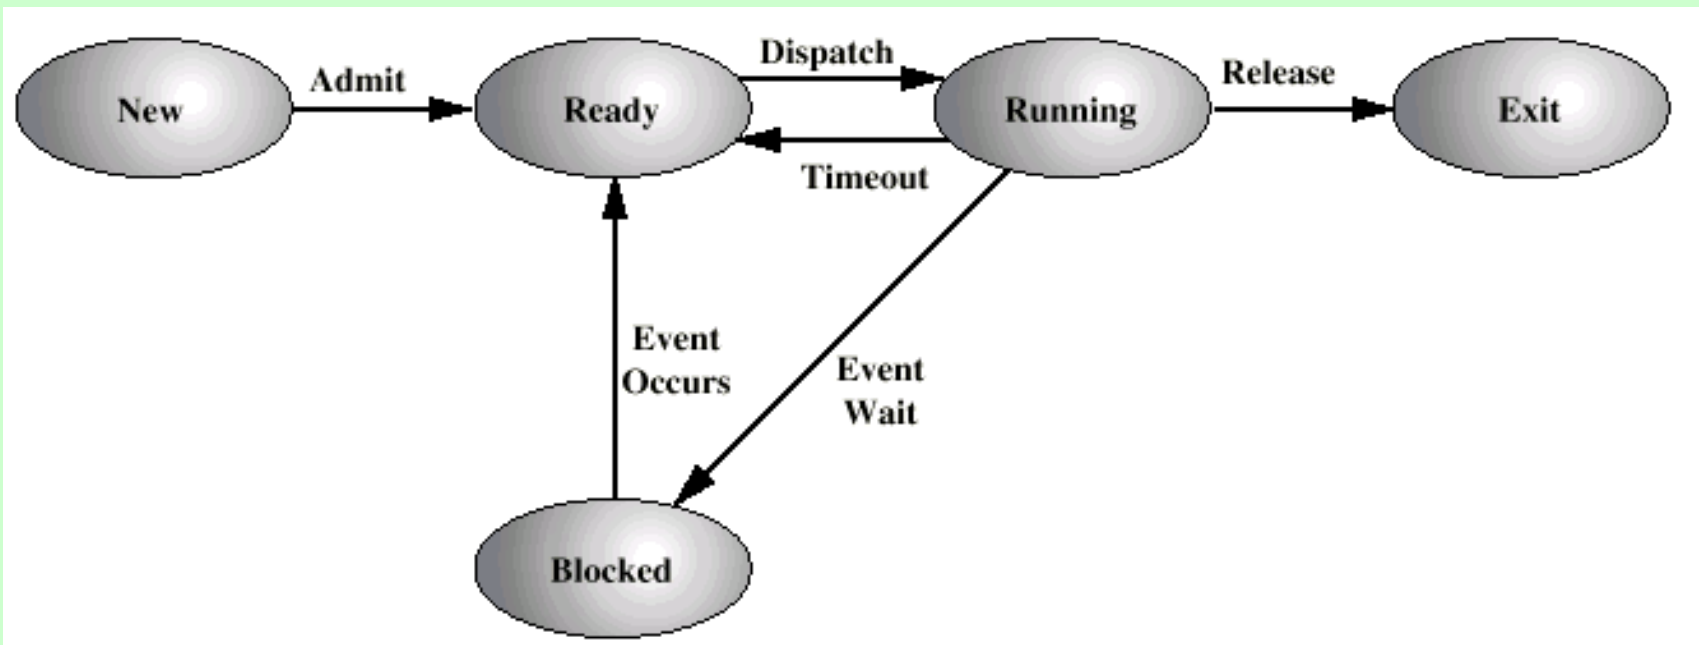
Medium Term Scheduling

- Part of the swapping function (later...)
- Usually based on the need to manage multi-programming
- If no virtual memory, memory management is also an issue

Short Term Scheduler

- Dispatcher
- Fine grained decisions of which job to execute next
- i.e. which job actually gets to use the processor in the next time slot

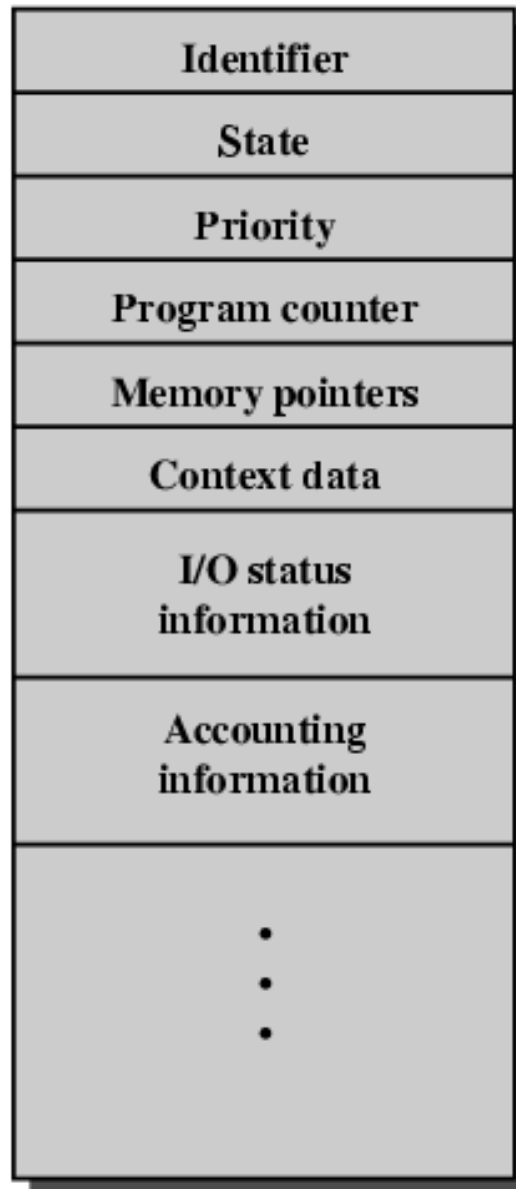
Five State Process Model




Process Control Block

- Identifier
- State
- Priority
- Program counter
- Memory pointers
- Context data
- I/O status
- Accounting information

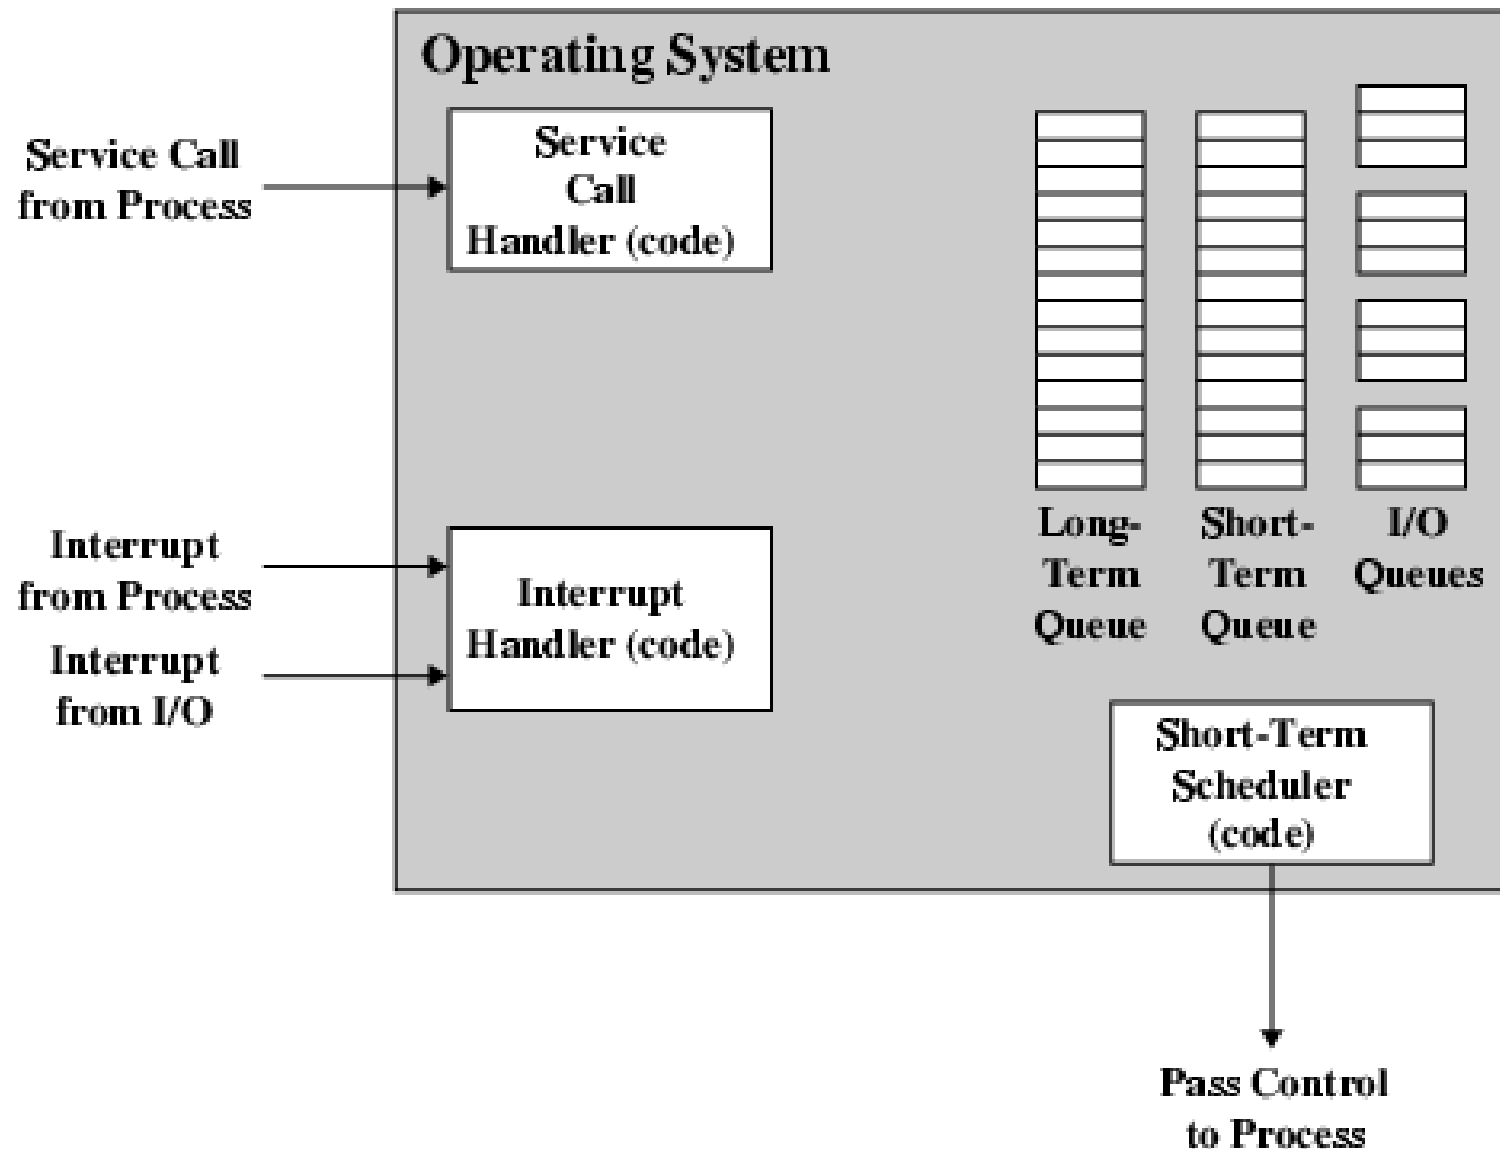
PCB Diagram



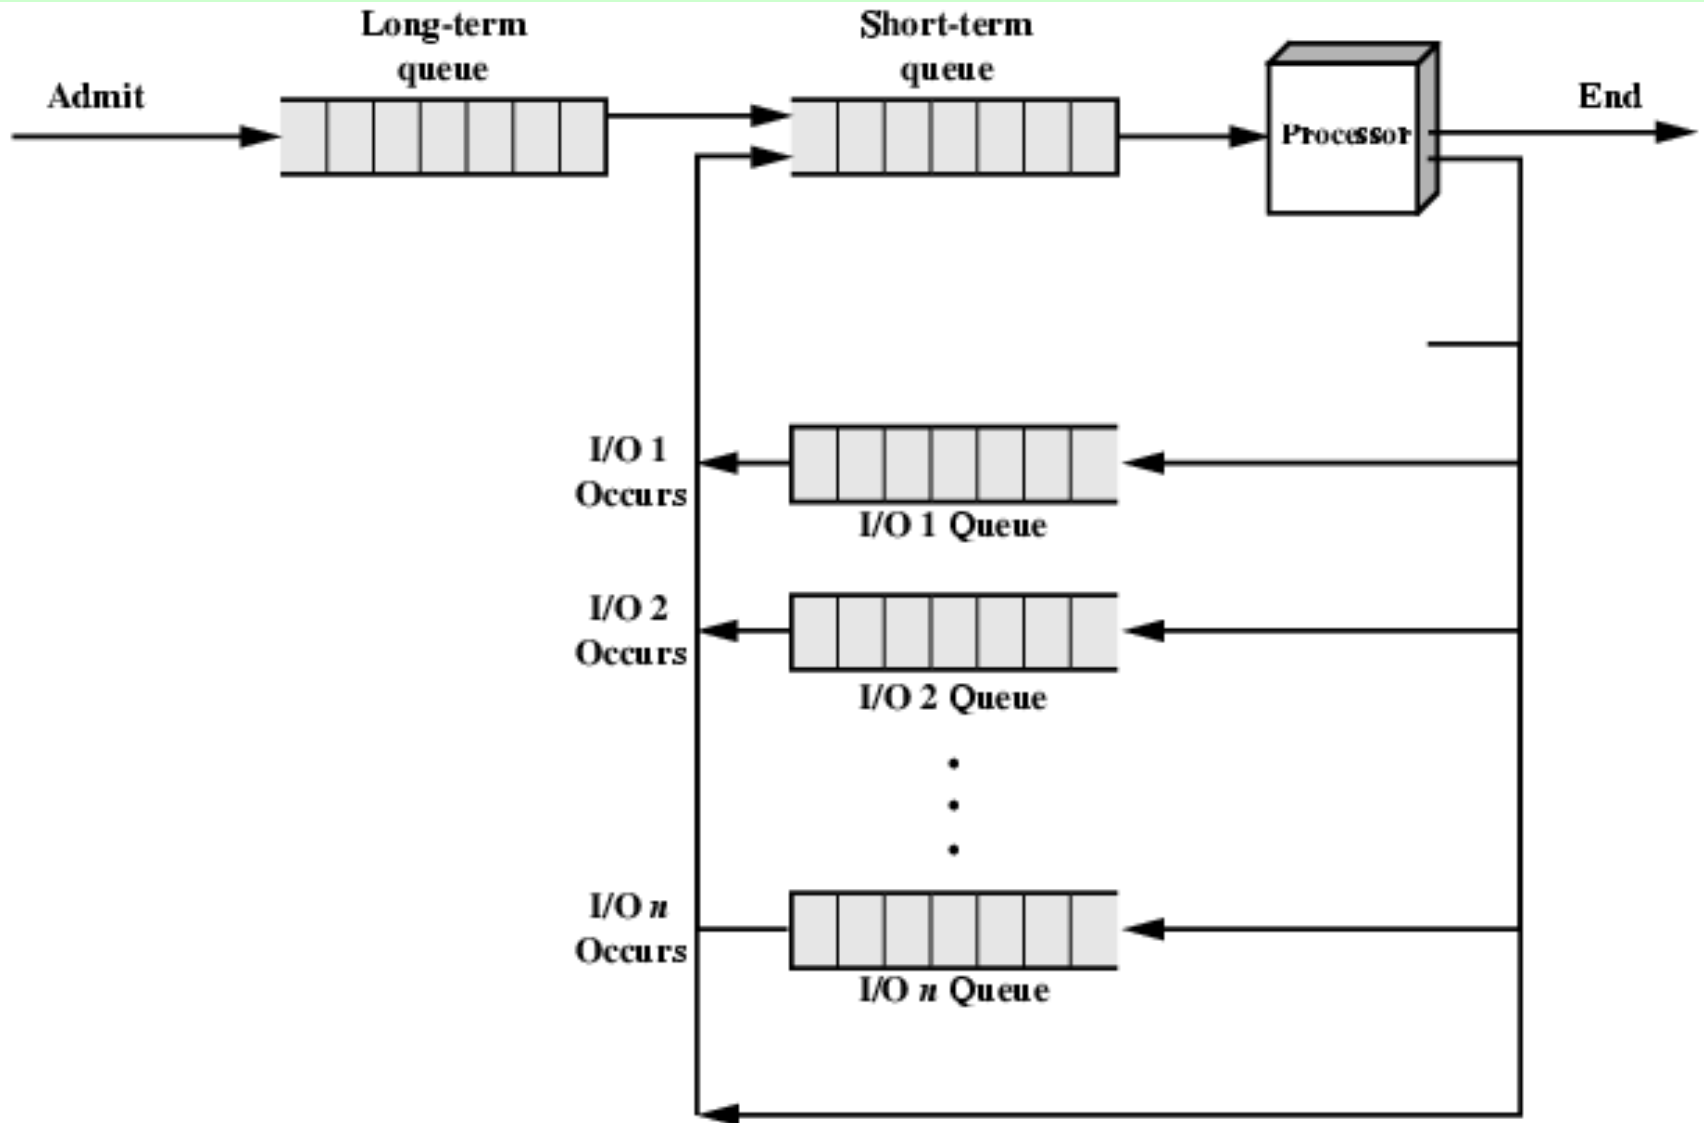
Scheduling Example



Key Elements of O/S



Process Scheduling



Memory Management

- Uni-program
 - Memory split into two
 - One for Operating System (monitor)
 - One for currently executing program
- Multi-program
 - “User” part is sub-divided and shared among active processes


Swapping

- Problem: I/O is so slow compared with CPU that even in multi-programming system, CPU can be idle most of the time
- Solutions:
 - Increase main memory
 - Expensive
 - Leads to larger programs
 - Swapping


What is Swapping?

- Long term queue of processes stored on disk
- Processes “swapped” in as space becomes available
- As a process completes it is moved out of main memory
- If none of the processes in memory are ready (i.e. all I/O blocked)
 - Swap out a blocked process to intermediate queue
 - Swap in a ready process or a new process
 - But swapping is an I/O process...

Use of Swapping



(a) Simple job scheduling




(b) Swapping


Partitioning

- Splitting memory into sections to allocate to processes (including Operating System)
- Fixed-sized partitions
 - May not be equal size
 - Process is fitted into smallest hole that will take it (best fit)
 - Some wasted memory
 - Leads to variable sized partitions

Fixed Partitioning



(a) Equal-size partitions



(b) Unequal-size partitions


Variable Sized Partitions (1)

- Allocate exactly the required memory to a process
- This leads to a hole at the end of memory, too small to use
 - Only one small hole - less waste
- When all processes are blocked, swap out a process and bring in another
- New process may be smaller than swapped out process
- Another hole

Variable Sized Partitions (2)

- Eventually have lots of holes (fragmentation)
- Solutions:
 - Coalesce - Join adjacent holes into one large hole
 - Compaction - From time to time go through memory and move all hole into one free block (c.f. disk de-fragmentation)

Effect of Dynamic Partitioning




Relocation

- No guarantee that process will load into the same place in memory
- Instructions contain addresses
 - Locations of data
 - Addresses for instructions (branching)
- Logical address - relative to beginning of program
- Physical address - actual location in memory (this time)
- Automatic conversion using base address

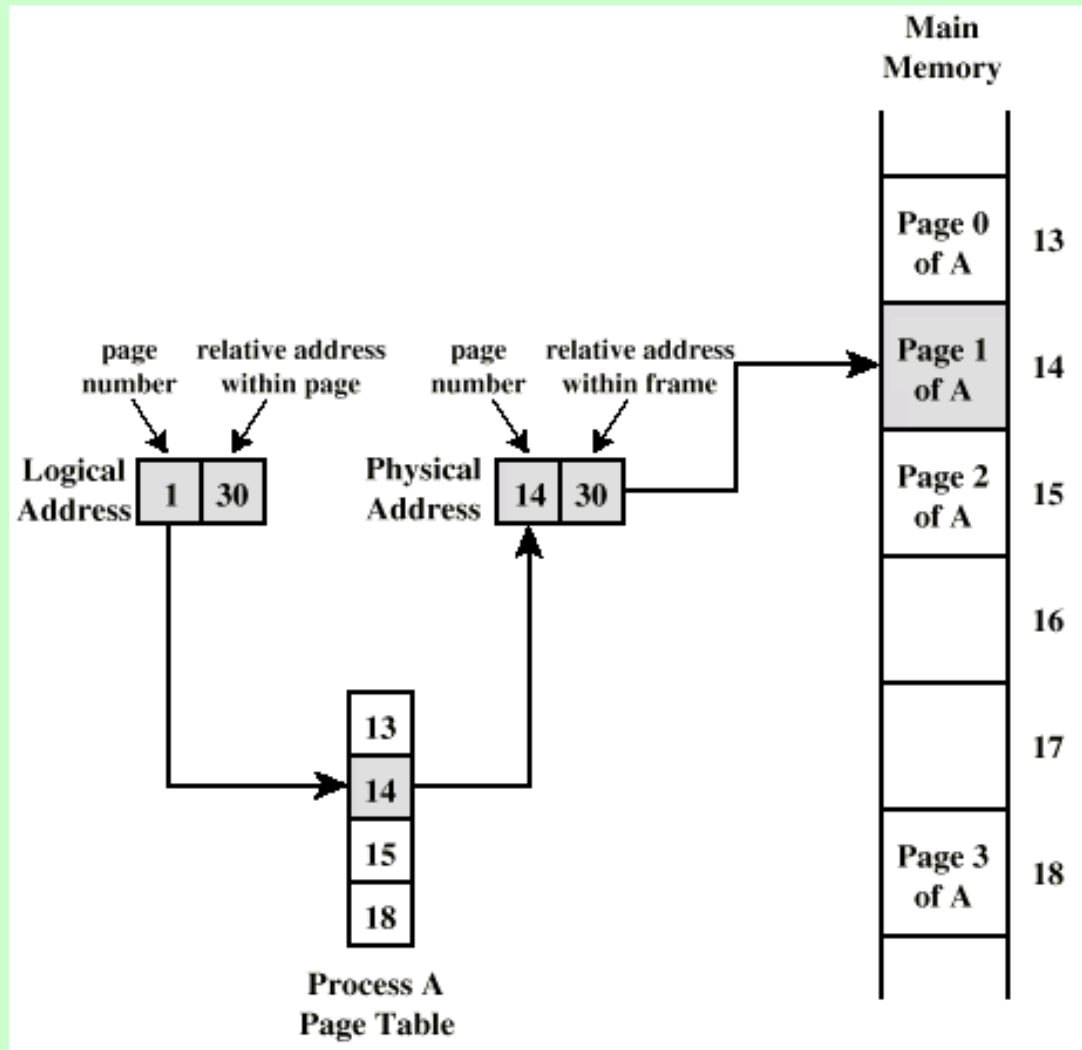
Paging

- Split memory into equal sized, small chunks -page frames
- Split programs (processes) into equal sized small chunks - pages
- Allocate the required number page frames to a process
- Operating System maintains list of free frames
- A process does not require contiguous page frames
- Use page table to keep track

Allocation of Free Frames



Logical and Physical Addresses - Paging



Virtual Memory

- Demand paging
 - Do not require all pages of a process in memory
 - Bring in pages as required
- Page fault
 - Required page is not in memory
 - Operating System must swap in required page
 - May need to swap out a page to make space
 - Select page to throw out based on recent history

Thrashing

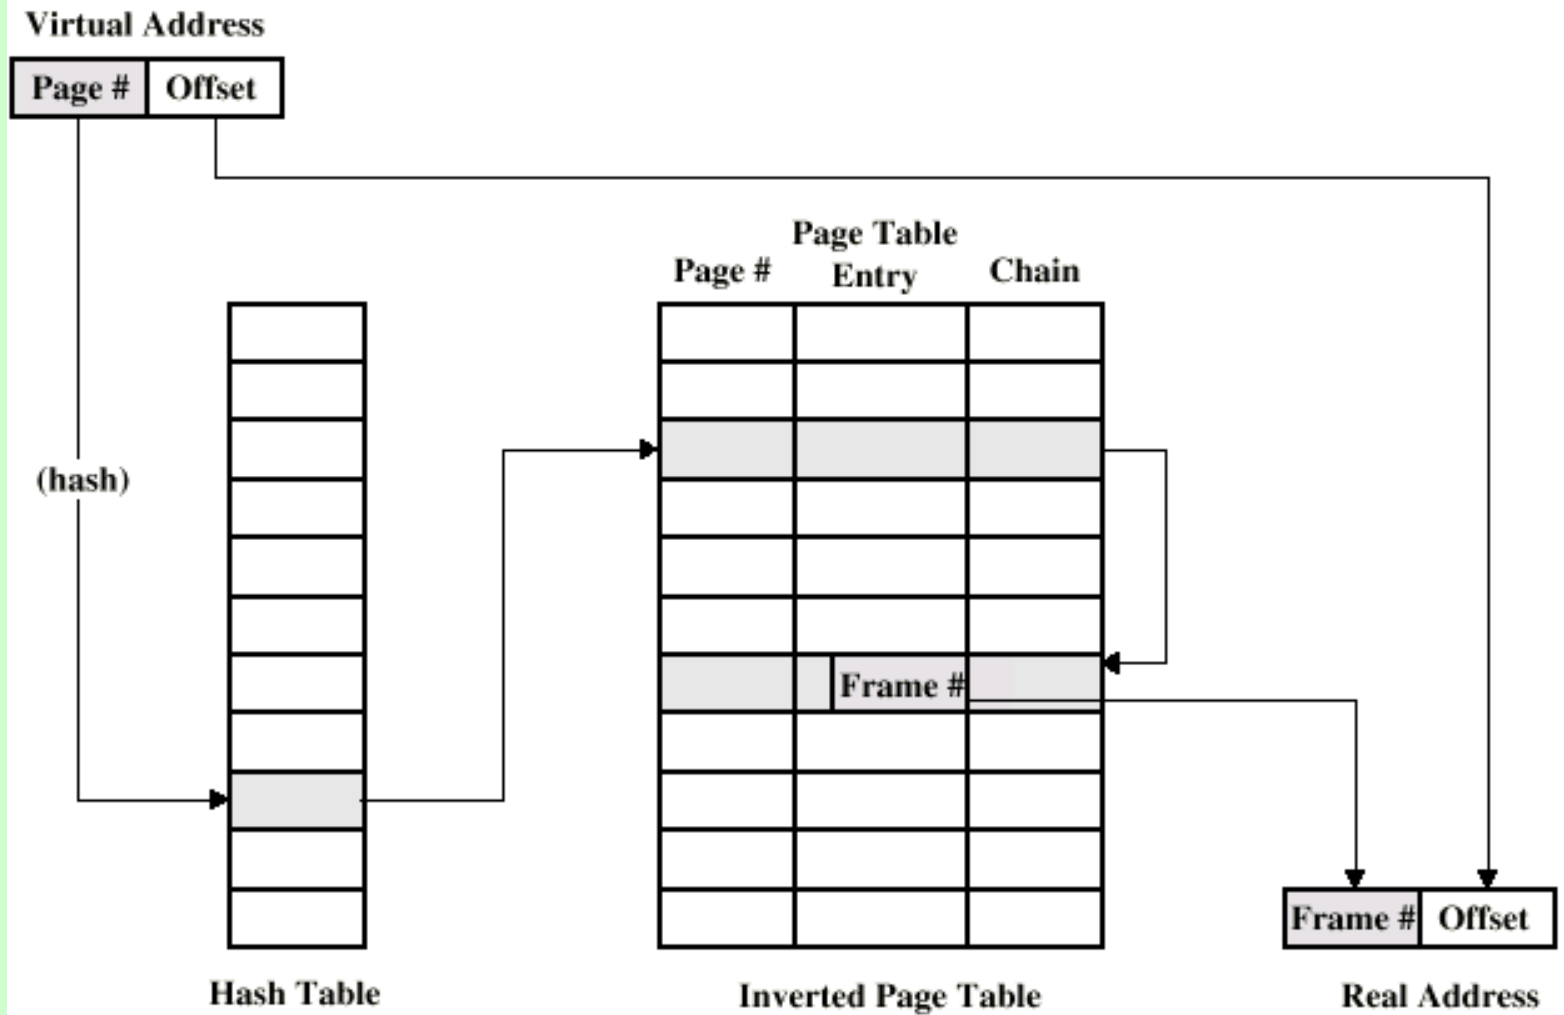
- Too many processes in too little memory
- Operating System spends all its time swapping
- Little or no real work is done
- Disk light is on all the time

- Solutions
 - Good page replacement algorithms
 - Reduce number of processes running
 - Fit more memory

Bonus

- We do not need all of a process in memory for it to run
 - We can swap in pages as required
 - So - we can now run processes that are bigger than total memory available!
-
- Main memory is called real memory
 - User/programmer sees much bigger memory - virtual memory

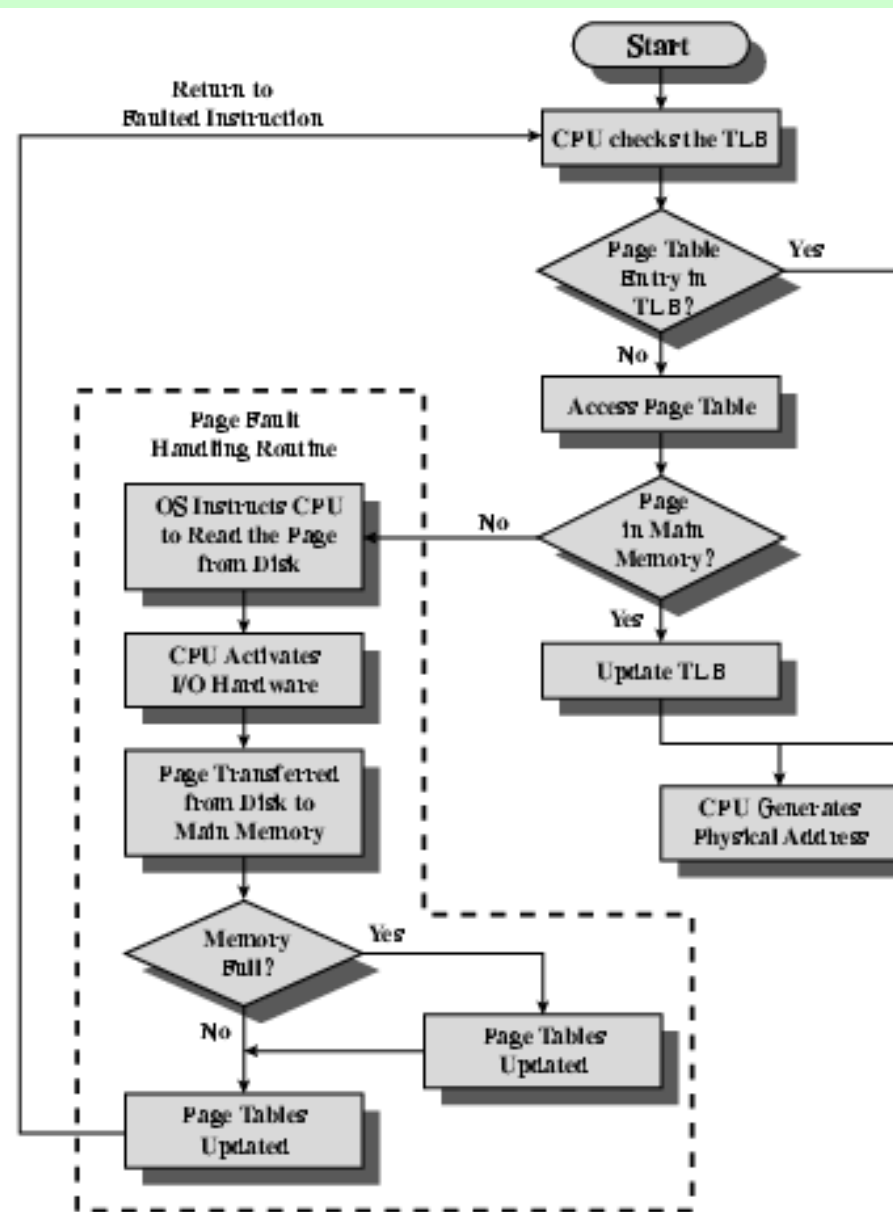
Inverted Page Table Structure



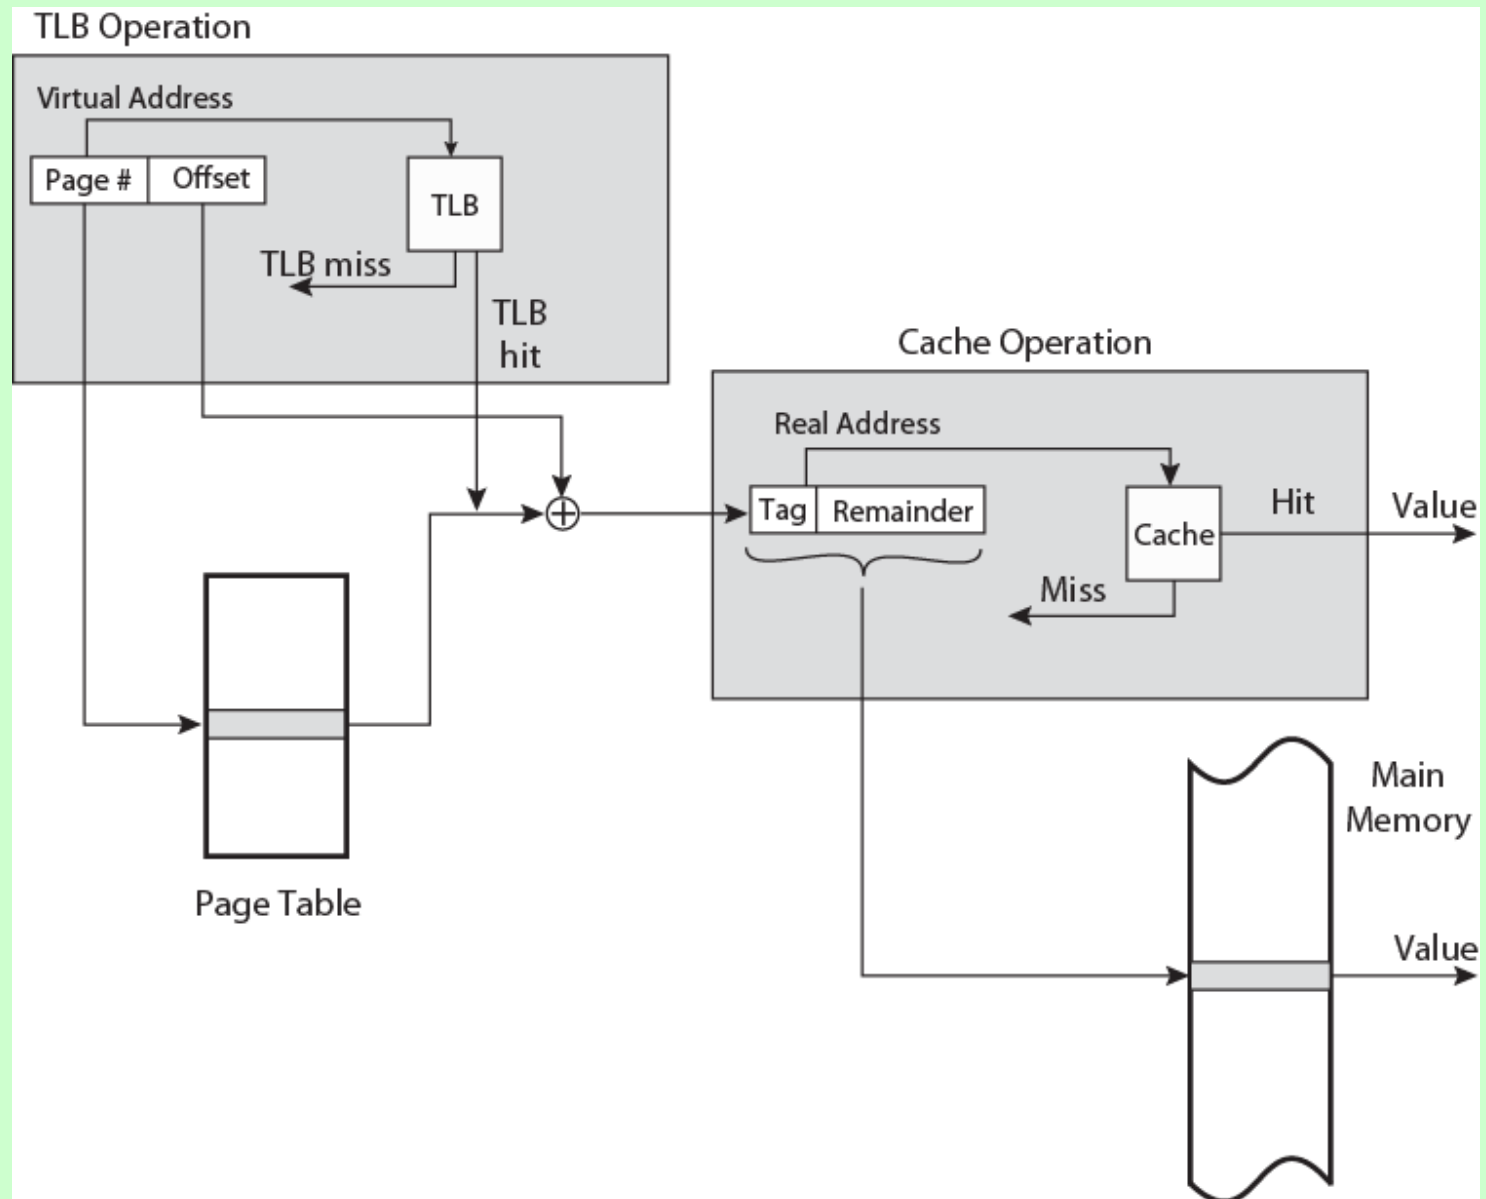
Translation Lookaside Buffer

- Every virtual memory reference causes two physical memory access
 - Fetch page table entry
 - Fetch data
- Use special cache for page table
 - TLB

TLB Operation



TLB and Cache Operation



Segmentation

- Paging is not (usually) visible to the programmer
- Segmentation is visible to the programmer
- Usually different segments allocated to program and data
- May be a number of program and data segments

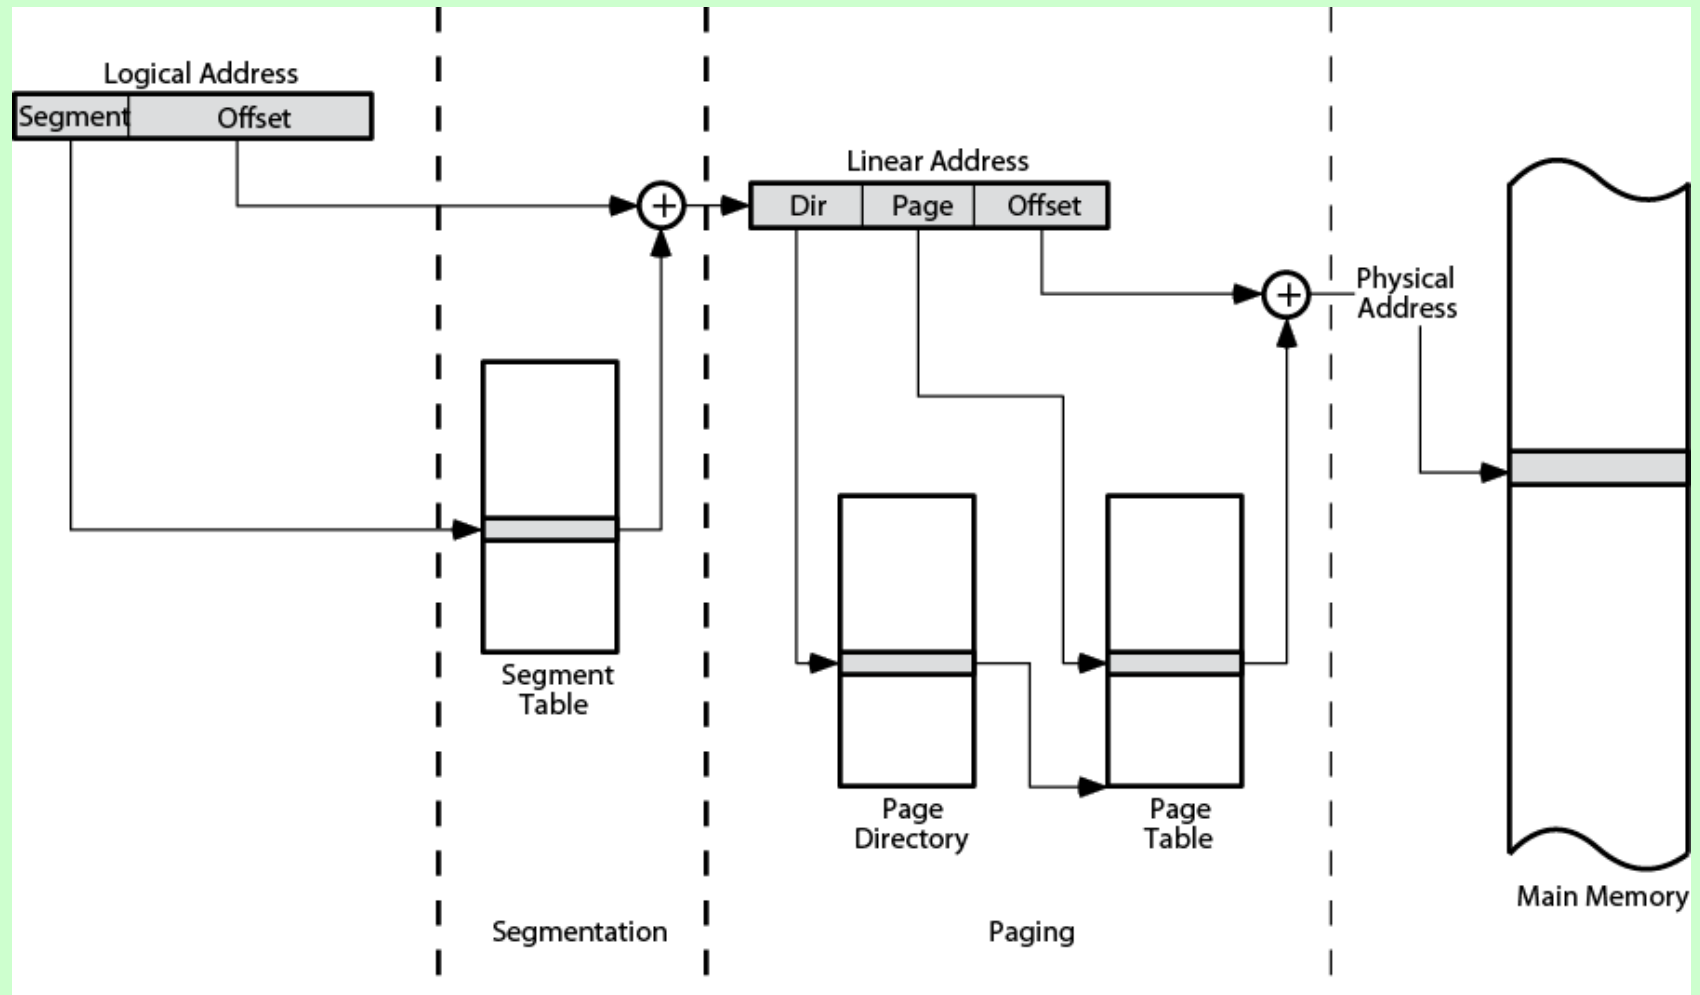
Advantages of Segmentation

- Simplifies handling of growing data structures
- Allows programs to be altered and recompiled independently, without re-linking and re-loading
- Lends itself to sharing among processes
- Lends itself to protection
- Some systems combine segmentation with paging

Pentium II

- Hardware for segmentation and paging
- Unsegmented unpaged
 - virtual address = physical address
 - Low complexity
 - High performance
- Unsegmented paged
 - Memory viewed as paged linear address space
 - Protection and management via paging
 - Berkeley UNIX
- Segmented unpaged
 - Collection of local address spaces
 - Protection to single byte level
 - Translation table needed is on chip when segment is in memory
- Segmented paged
 - Segmentation used to define logical memory partitions subject to access control
 - Paging manages allocation of memory within partitions
 - Unix System V

Pentium II Address Translation Mechanism



Pentium II Segmentation

- Each virtual address is 16-bit segment and 32-bit offset
- 2 bits of segment are protection mechanism
- 14 bits specify segment
- Unsegmented virtual memory $2^{32} = 4\text{Gbytes}$
- Segmented $2^{46} = 64\text{ terabytes}$
 - Can be larger – depends on which process is active
 - Half (8K segments of 4Gbytes) is global
 - Half is local and distinct for each process

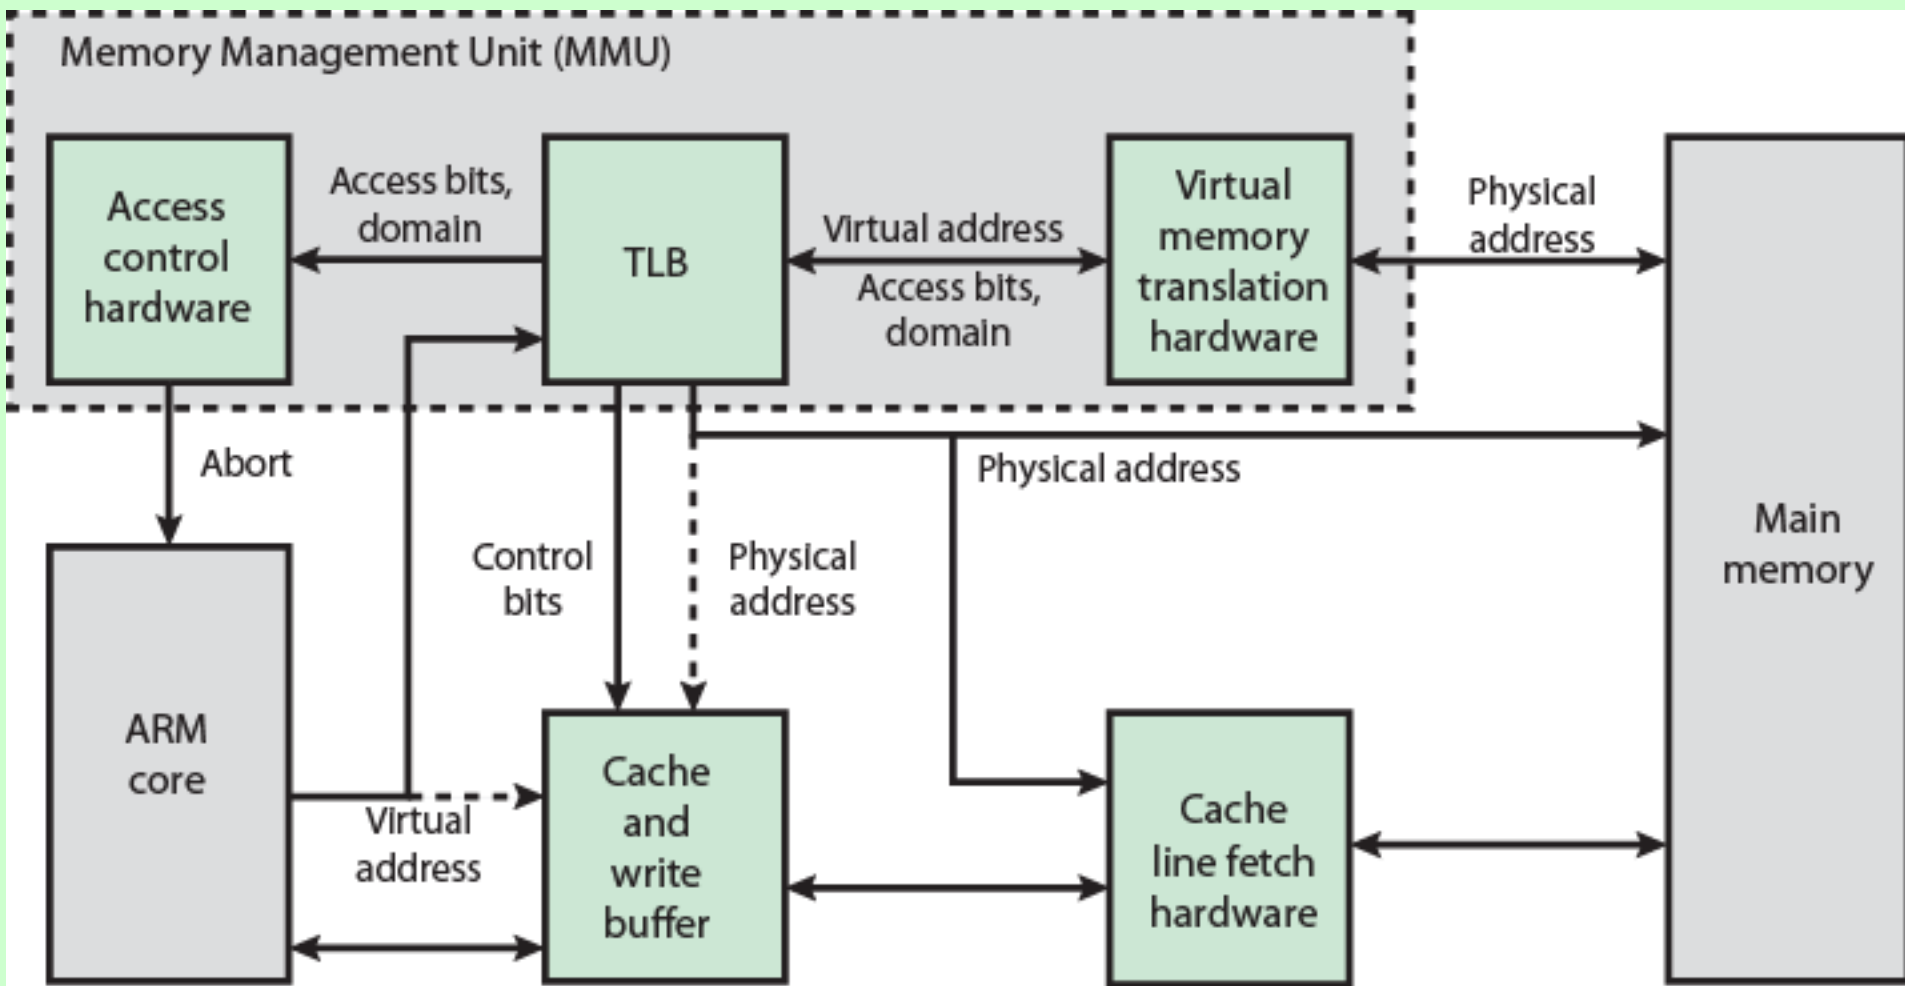
Pentium II Protection

- Protection bits give 4 levels of privilege
 - 0 most protected, 3 least
 - Use of levels software dependent
 - Usually level 3 for applications, level 1 for O/S and level 0 for kernel (level 2 not used)
 - Level 2 may be used for apps that have internal security e.g. database
 - Some instructions only work in level 0

Pentium II Paging

- Segmentation may be disabled
 - In which case linear address space is used
- Two level page table lookup
 - First, page directory
 - 1024 entries max
 - Splits 4G linear memory into 1024 page groups of 4Mbyte
 - Each page table has 1024 entries corresponding to 4Kbyte pages
 - Can use one page directory for all processes, one per process or mixture
 - Page directory for current process always in memory
 - Use TLB holding 32 page table entries
 - Two page sizes available 4k or 4M

ARM Memory System Overview



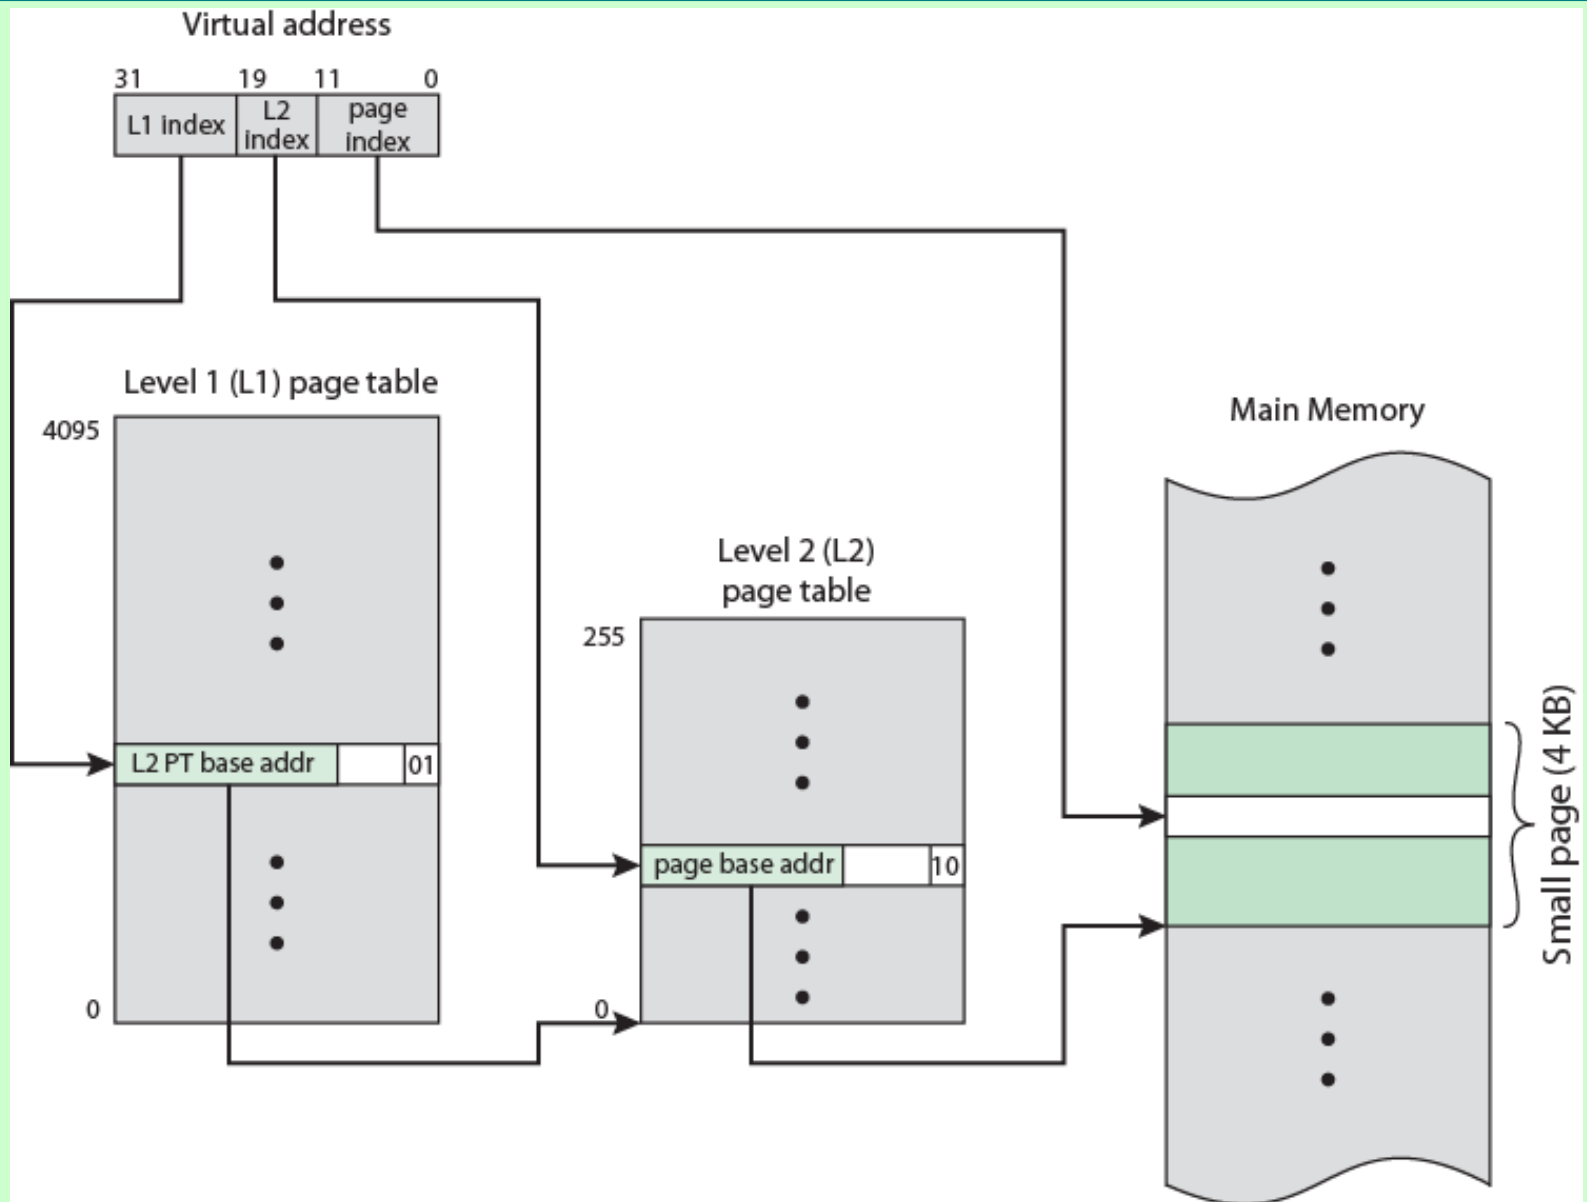
ARM Memory Management

- Virtual memory translation
 - One or two levels of tables
- Translation lookaside buffer (TLB)
 - Cache of recent entries
 - If available, TLB directly sends physical address to main memory
- Data exchanged between processor and main memory via cache
- Logical cache organization
 - On cache miss, ARM supplies address directly to cache as well as TLB
- Physical cache organization
 - TLB must supply physical address to cache
- Access control bits also in translation tables

Virtual Memory Address Translation

- Memory access based on either sections or pages
- Supersections (optional)
 - 16-MB blocks of main memory
- Sections
 - 1-MB blocks of main memory
- Large pages
 - 64-KB blocks of main memory
- Small pages
 - 4-KB blocks of main memory
- Sections and supersections allow mapping of large region of memory with single TLB entry
- Additional access control mechanisms
 - Small pages use 1KB subpages
 - Large pages use 16KB subpages
- Two level translation table held in main memory
 - First-level table holds section and supersection translations, and pointers to second-level tables
 - Second-level tables: Hold both large and small page translations
- MMU
 - Translates virtual to physical addresses
 - Derives and checks access permission
 - After TLB miss
- Start with first-level fetch
 - Section-mapped access only requires first-level fetch
 - Page-mapped access requires second-level fetch

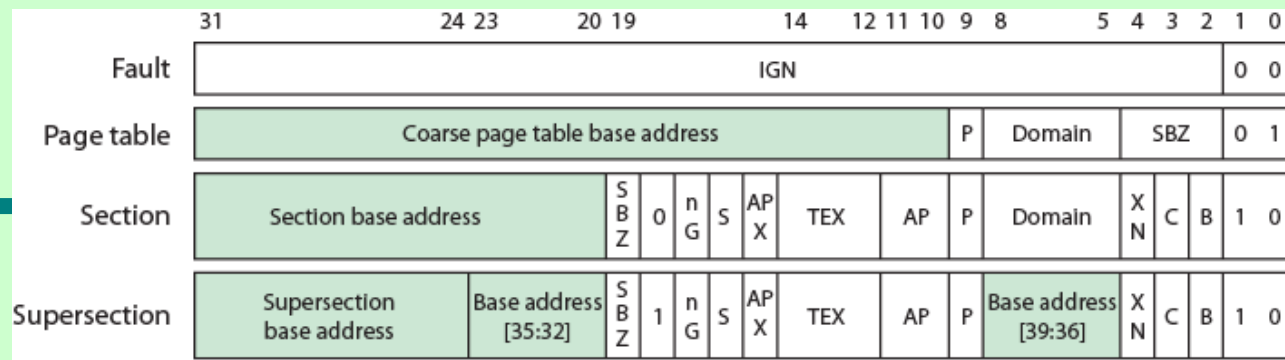
ARM Virtual Memory Address Translation for Small Pages - Diagram



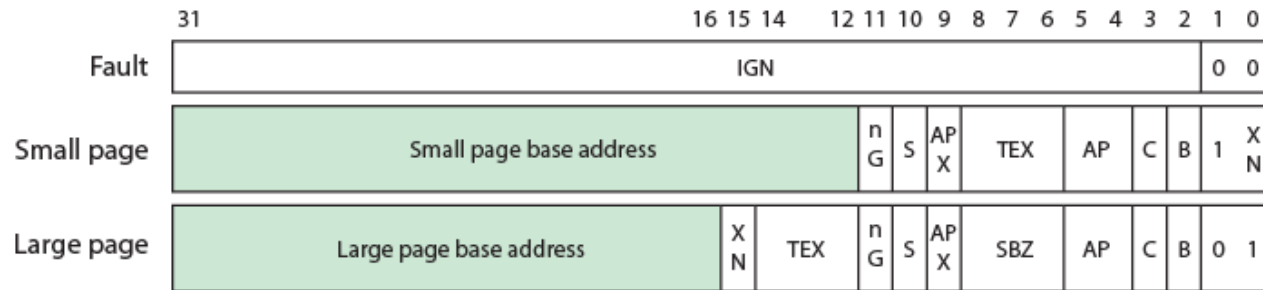
ARM Virtual Memory Address Translation for Small Pages

- Single L1 page table
 - 4K 32-bit entries
 - Each L1 entry points to L2 page table
- Each L2 page table
 - 255 32-bit entries
 - Each L2 entry points to 4-KB page in main memory
- 32-bit virtual address
 - Most significant 12 bits index L1 page table
 - Next 8 bits index relevant L2 page table
 - Least significant 12 bits index a byte in relevant main memory page
- Similar procedure for large pages
- Sections and supersection only use L1 page table


ARMv6 Memory Management Formats



(a) Alternative first-level descriptor formats



(b) Alternative second-level descriptor formats



(c) Virtual memory address formats

ARM Memory-Management Parameters

- **Access Permission (AP), Access Permission Extension (APX)**
 - Control access to corresponding memory region
 - Access without required permissions raises Permission Fault
- **Bufferable (B) bit**
 - With TEX bits, how write buffer is used
- **Cacheable (C) bit**
- Can memory region be mapped through cache?
- **Domain**
 - Collection of memory regions
 - Access control can be applied on the basis of domain
- **not Global (nG)**
 - Translation marked as global (0), or process specific (1)?
- **Shared (S)**
 - Translation is for not-shared (0), or shared (1) memory?
- **SBZ**
 - Should be zero
- **Type Extension (TEX)**
 - Together with B and C bits, control accesses to caches
 - How write buffer is used
 - If memory region is shareable
 - Must be kept coherent
- **Execute Never (XN)**
 - Region is executable (0) or not executable (1)?

Memory Management Formats – L1 table

- L1 table
 - Entry describes how associated 1-MB virtual address range is mapped
- Bits [1:0] = 00
 - Virtual addresses unmapped
 - Attempts to access generate translation fault
- Bits [1:0] = 01
 - Physical address of L2 page table which specifies how associated virtual address range is mapped
- Bits [1:0] = 01 and bit 19 = 0
 - Section descriptor
 - Supersection descriptor
- Entries with bits [1:0] = 11
 - Reserved

L2 Table

Small and Large Pages

- For memory structured into pages
- L1 page entry bits [31:10] point to a L2 page table
- Small pages
 - L2 entry holds 20-bit pointer to base address of 4-KB page in main memory
- Large pages
 - Virtual address includes 12-bit index to L1 table and an 8-bit index to L2 table
 - 64-KB large pages have 16 bit page index portion
 - Four-bit overlap between page index field and L2 table index field
 - Page table entries in L2 page table replicated 16 times
 - L2 page table reduced from 256 entries to 16 if all refer to large pages
 - L2 page can service mixed large and small pages

L2 Table

Sections and Supersections

- Sections or supersections
 - One-level page table access
- Sections
 - L1 entry Bits [31:20] hold 12-bit pointer to 1-MB section
- For supersections
 - L1 bits [31:24] hold 8-bit pointer to base of the 16-MB section
- Page table entry replication is required
 - Supersections L1 table index portion of virtual address overlaps 4 bits with supersection index portion of virtual address
 - 16 identical L1 page table entries
- Physical address space can be expanded by up to eight additional address bits
 - Bits [23:20] and [8:5]
 - Implementation dependent
 - Interpreted as extending physical memory by up to $2^8 = 256$
 - Physical memory may be up to 256 times memory space available to each individual process

Access Control

- Region of memory can be designated as no access, read only, or read/write
- Region can be designated privileged access (operating Systems) only
- Domain
 - Collection of sections and/or pages with particular access permissions
 - 16
 - Multiple processes can use same translation tables while maintaining some protection from each other
 - Page table entry and TLB entry contain domain field
 - Two-bit field in the Domain Access Control Register controls access to each domain
 - Whole memory areas can be swapped very efficiently
- Clients
 - Must observe permissions of individual sections and/or pages in domain
- Managers
 - Control domain behavior
 - Sections and pages in domain access
 - Bypass access permissions for table entries in domain
- Programs can be
 - Client of some domains
 - Manager of other domains
 - No access to remaining domains

Required Reading

- Stallings chapter 8
- Stallings, W. [2004] *Operating Systems*, Pearson
- Loads of Web sites on Operating Systems